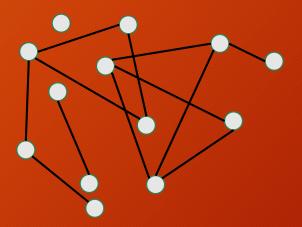
Parallel Connected Components

The Parallelepipeds:
Fabian Meier
Gustavo Segovia
Seraiah Walter

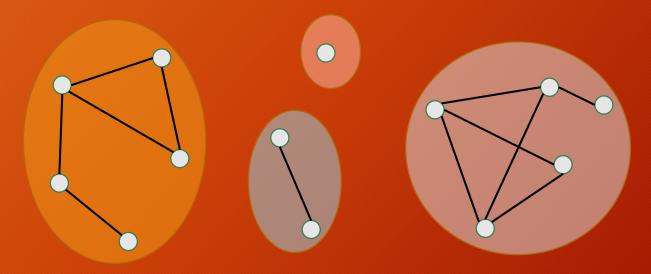
Our goal

- Develop a high performance algorithm for connected components
- Graph fits in memory



Our goal

- Develop a high performance algorithm for connected components
- Graph fits in memory



Algorithms - Related work

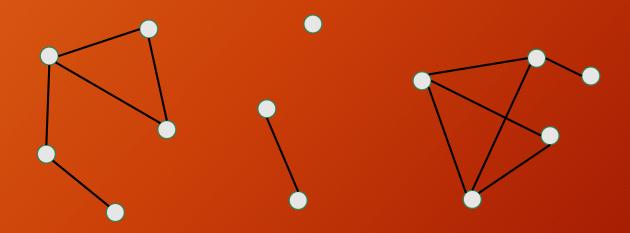
- Basic serial
 - Traversal (BFS/DFS) O(n + m)
 - Union find O(m log(n))
- Boost
 - Serial
 - Parallel
- Randomized Contraction Parallel Connected Components
 - http://www3.cs.stonybrook.edu/~rezaul/Spring-2012/CSE613/CSE613-lecture-11.pdf

Algorithms - Ours

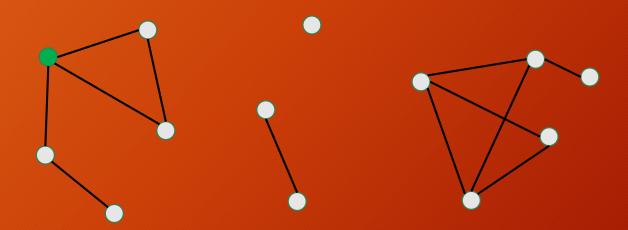
- Lock free parallel traversal
 - Parallel BFS
 - Parallel BFS with atomics
- Parallel spanning tree inspired by union find

Algorithm details

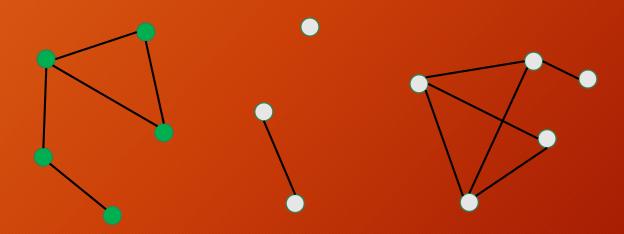
Unmarked vertices



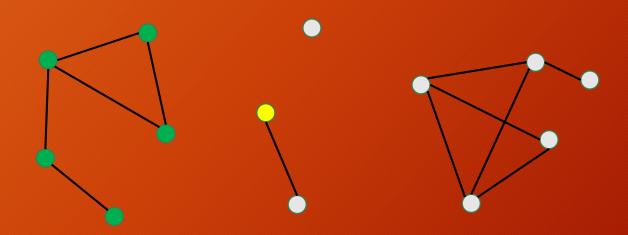
Algorithm starts at one vertex



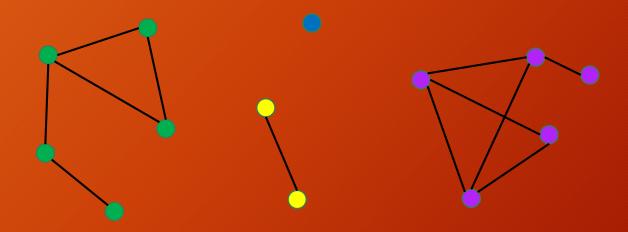
And transverses all connected vertices

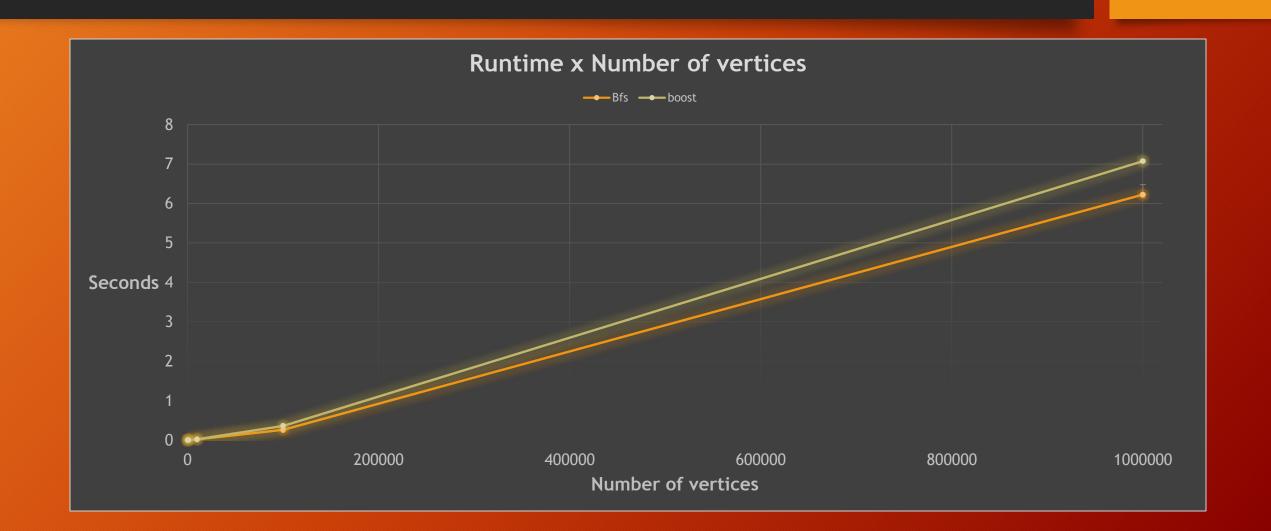


Continues to next unmarked vertex



And marks each component





Union find

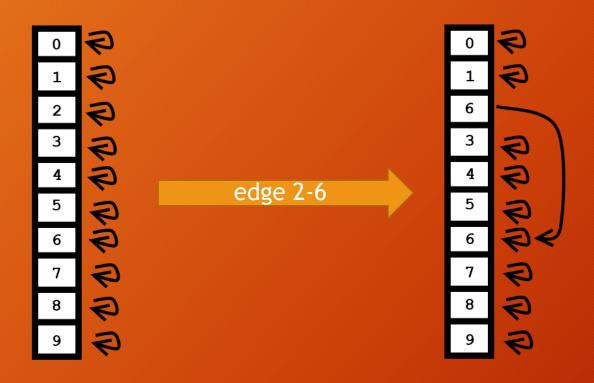
Union Find



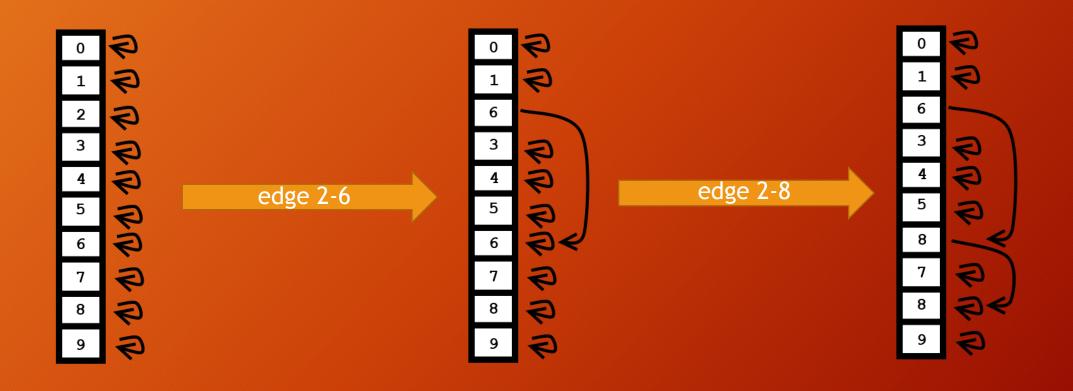


F

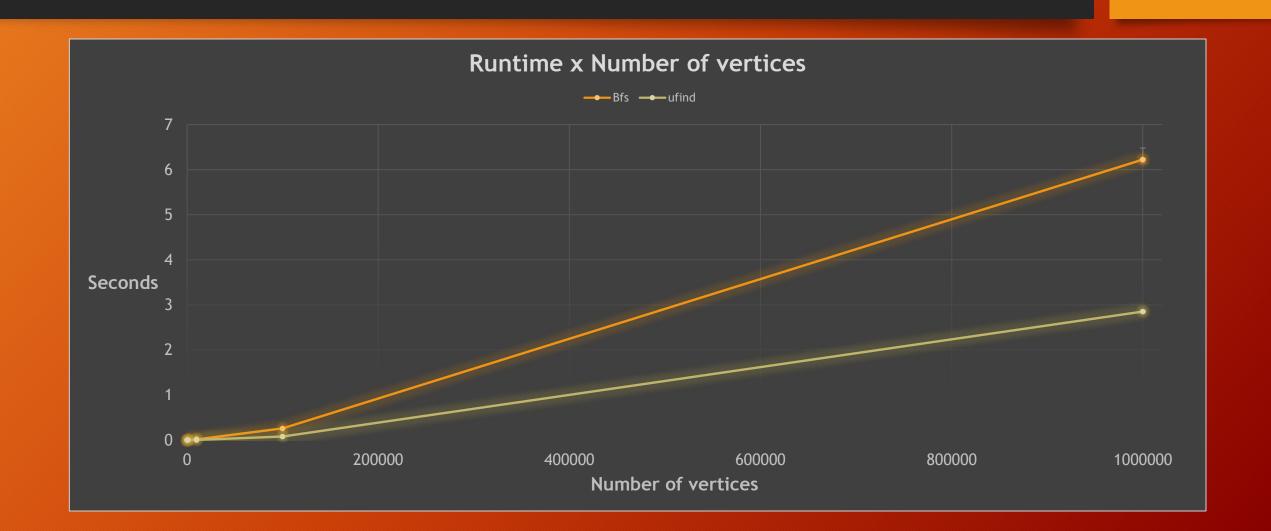
Union Find

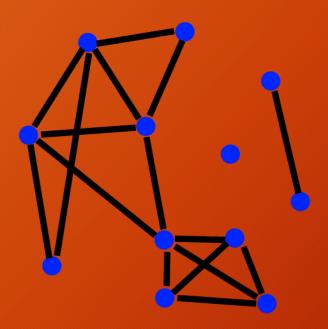


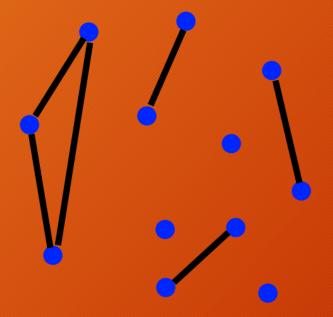
Union Find

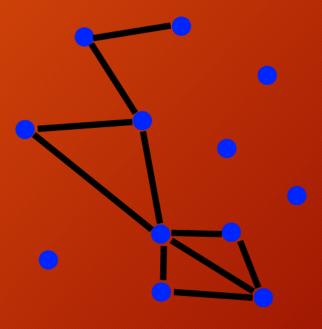


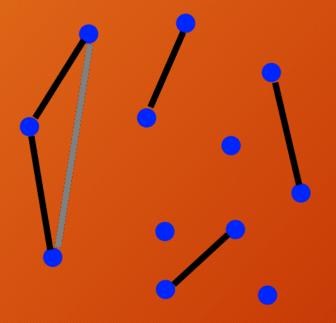
Union find

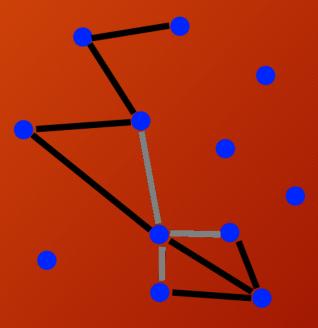


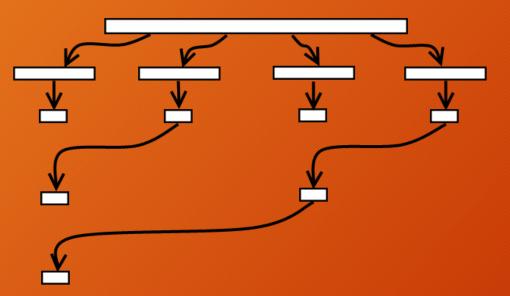




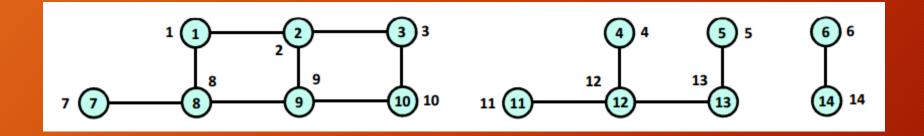




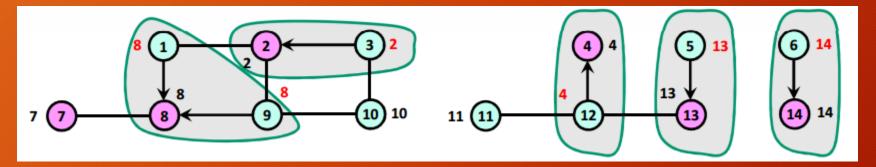




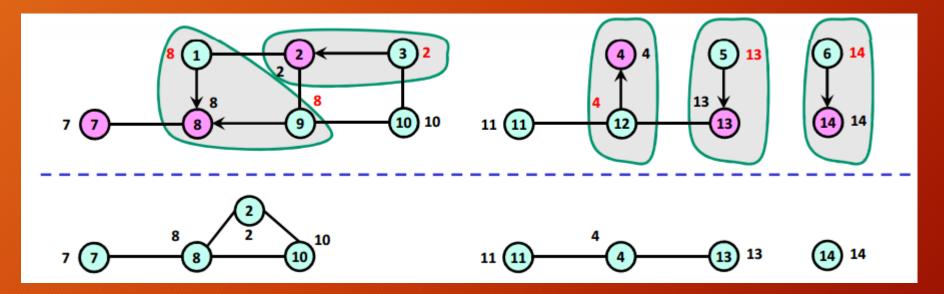
- Runtime: O(m/p+ n*log(p))
- Strong scaling in number of processors



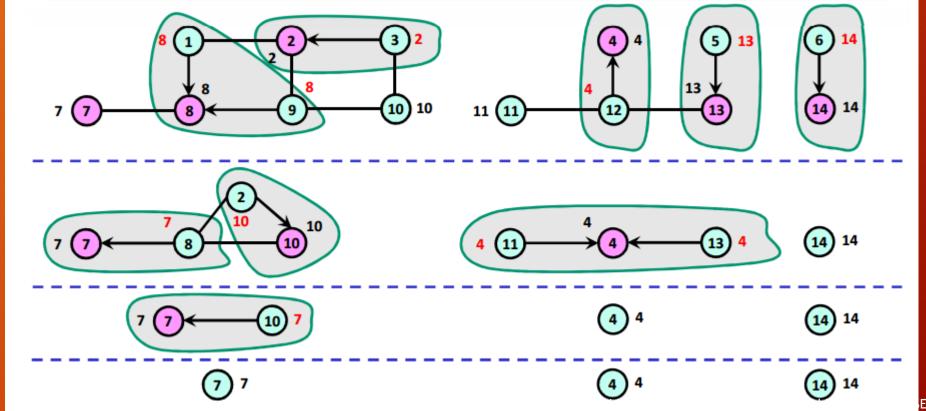
• Each vertex is given a random color (pink / light blue)



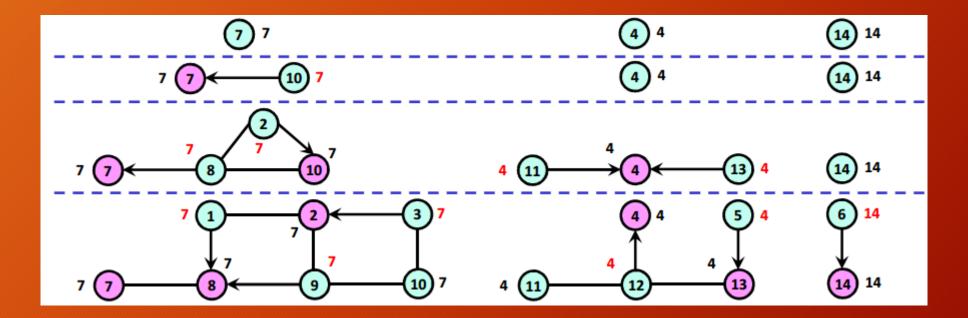
• Each contraction leads to less edges

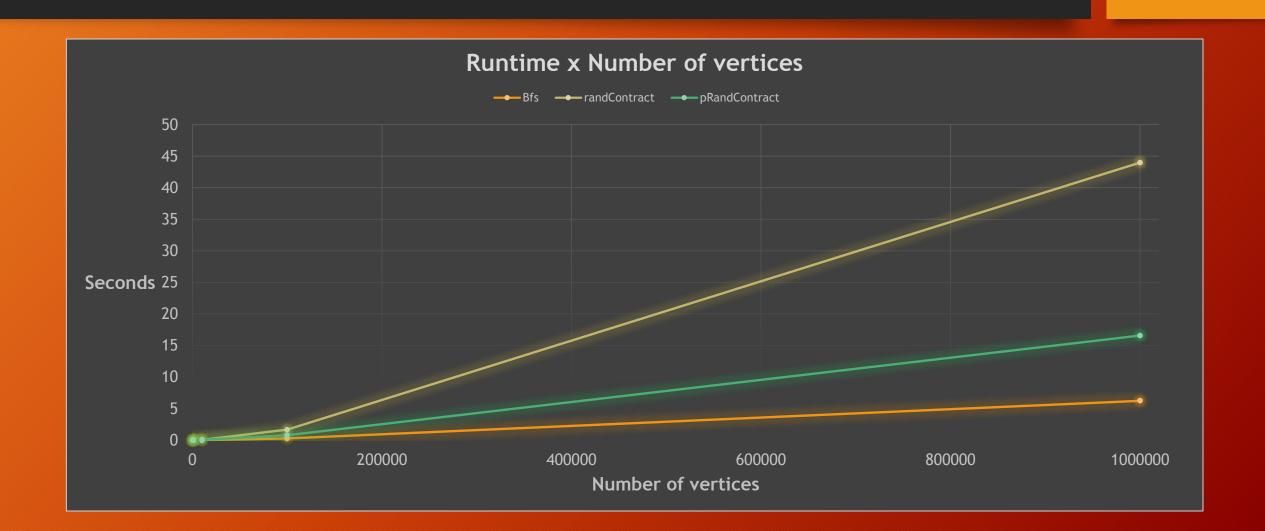


Contract until no edges are left

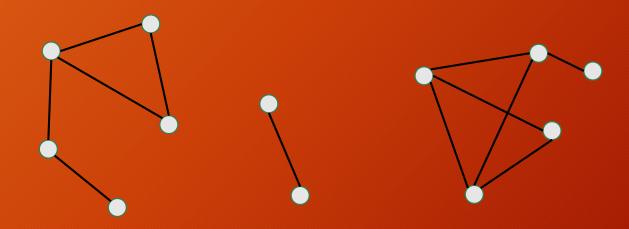


Reverse the path to find the vertices component number

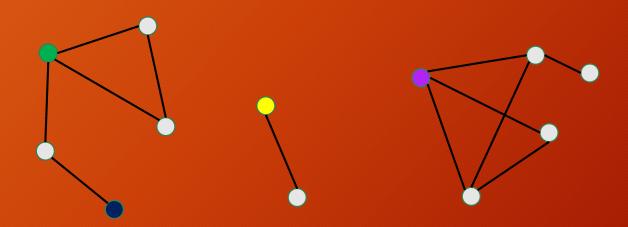




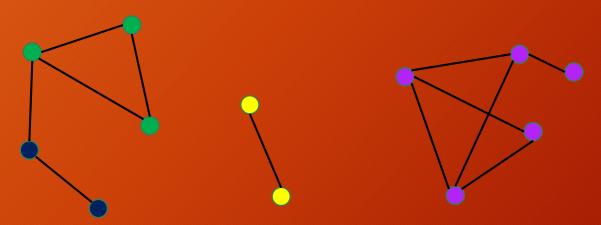
Unmarked vertices



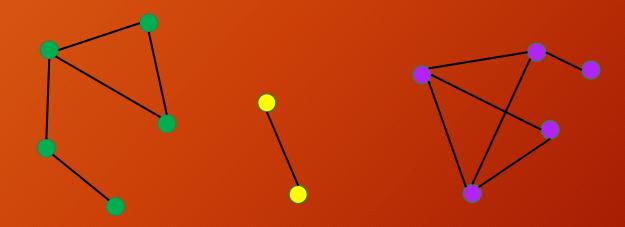
• Each thread starts at a vertex



- Vertices are marked
- When one thread hits a component marked by another thread, it makes an entry in the merge table



Vertices components are merged



 Problem, thread might not notice it is marking a vertex that was already marked by another thread

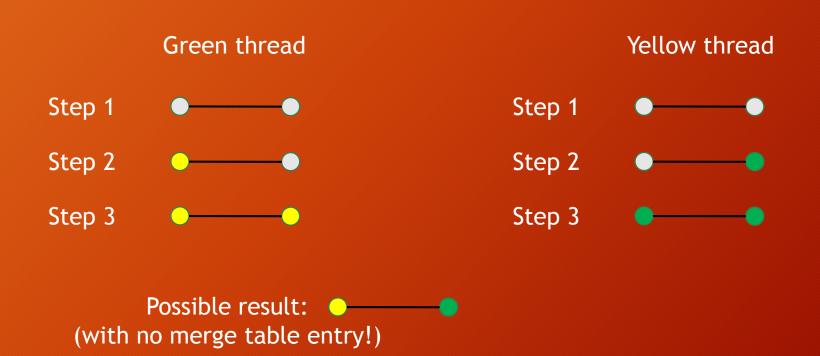
 Problem, thread might not notice it is marking a vertex that was already marked by another thread



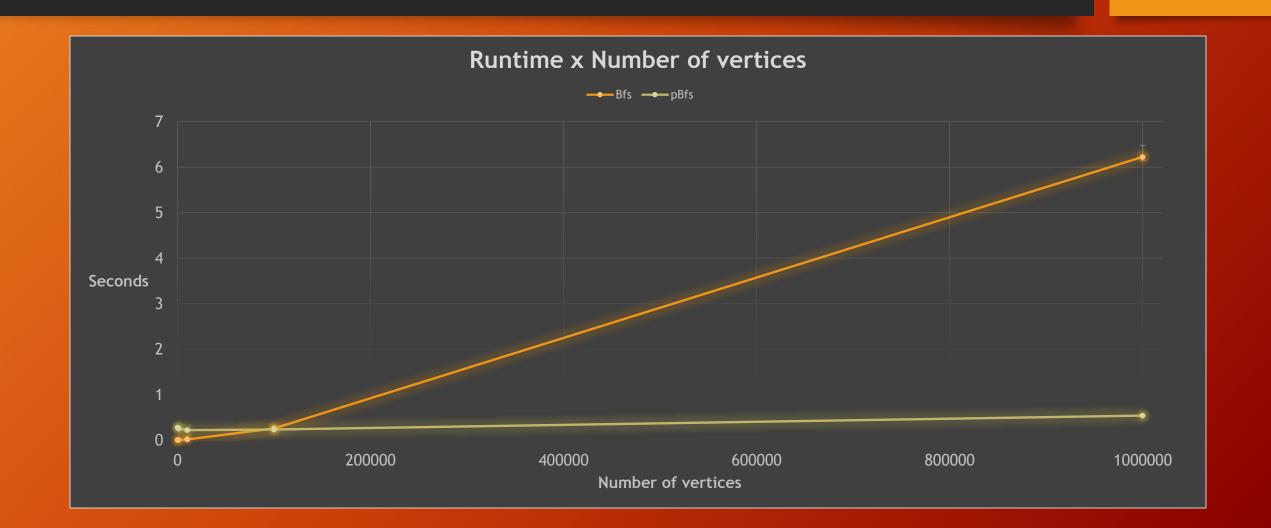








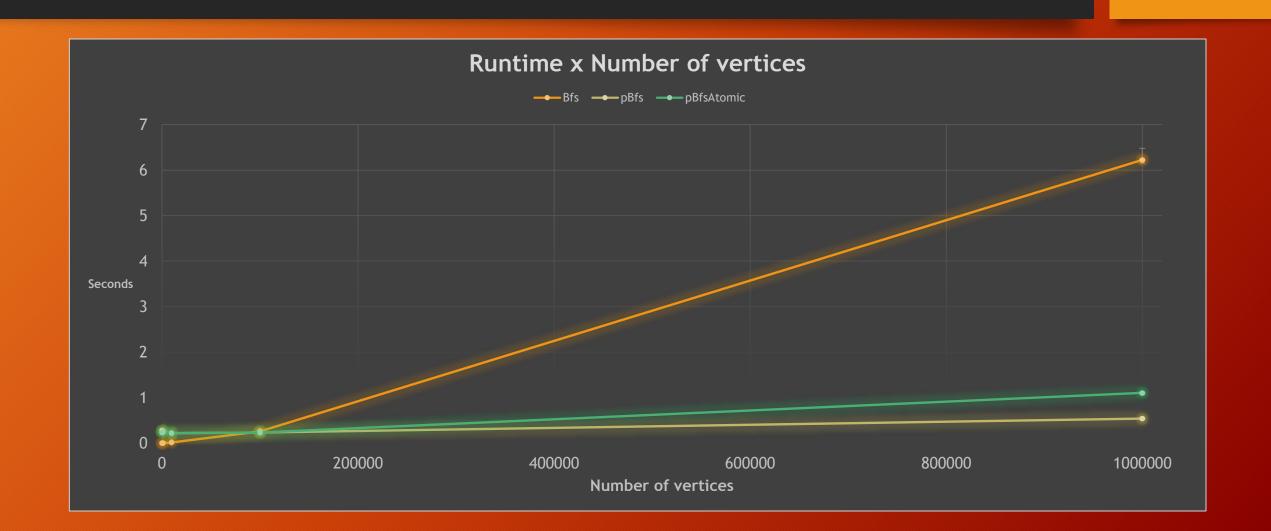
- Solution 1: check if algorithm output is correct. If not, run again with 1 thread
- Fast to check
- Problem is very unlikely. Slow execution is amortized



Parallel traversal with atomics

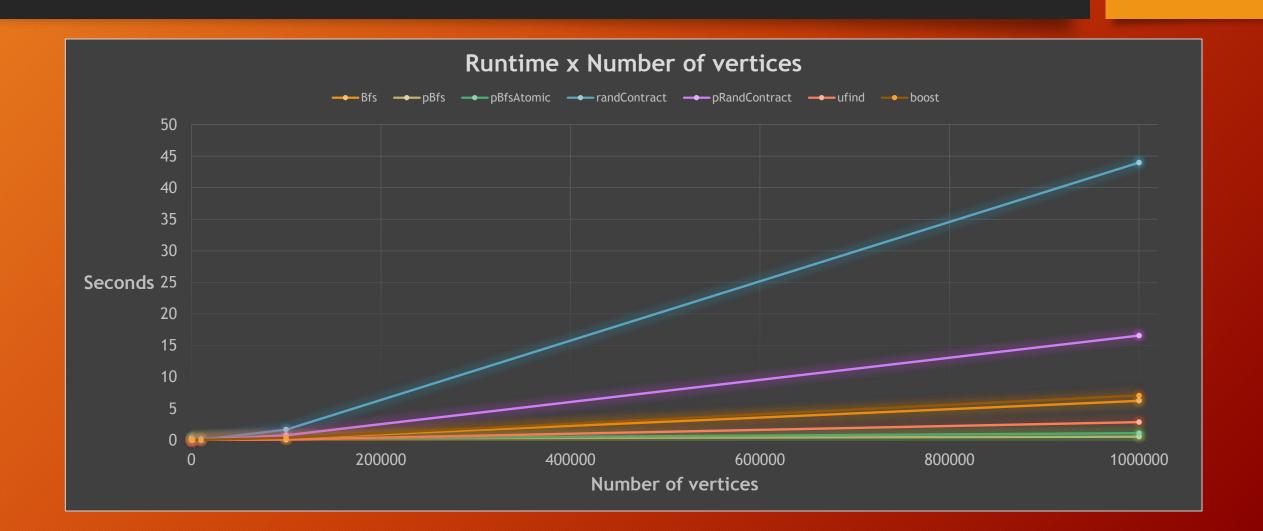
- Solution 2: Use test and set (atomic_flag)
- Every time a thread marks a vertex, it makes sure no other thread marked it before with test and set
- Each vertex has is own atomic_flag

Parallel traversal with atomics

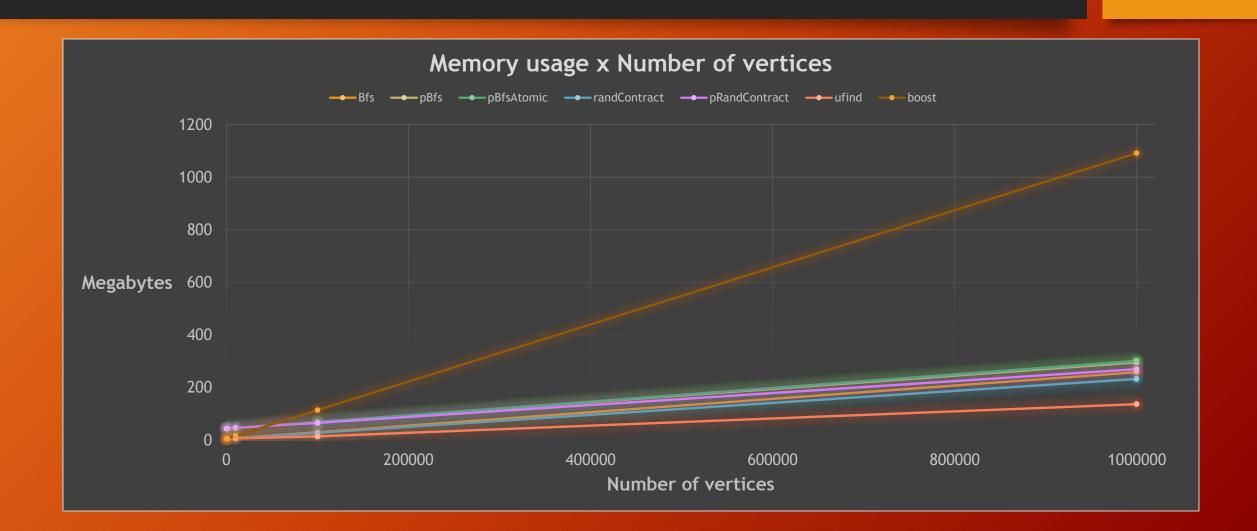


Algorithm results

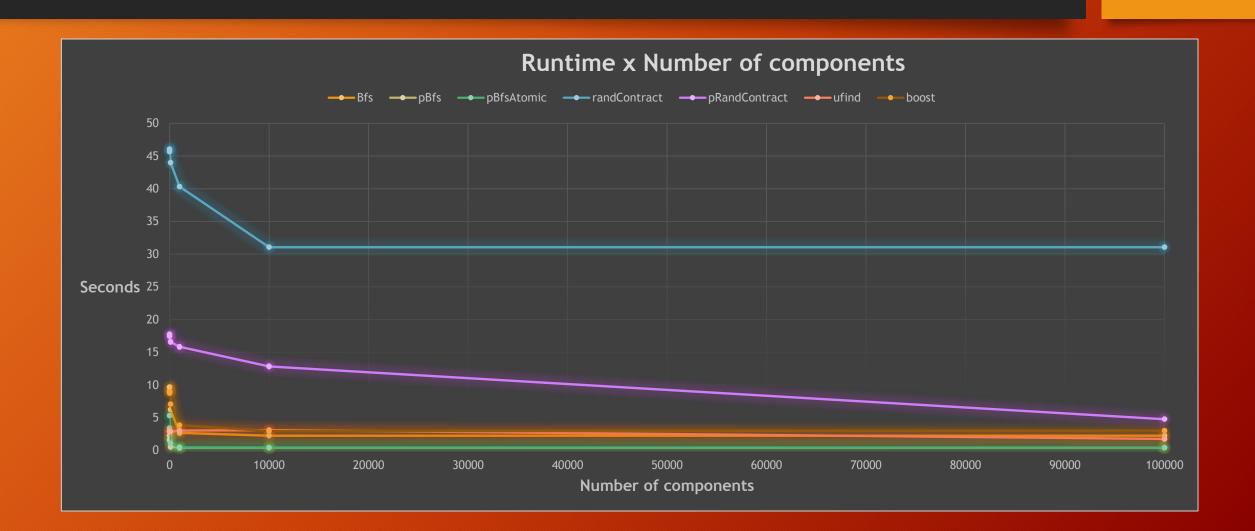
Runtime varying the number of vertices



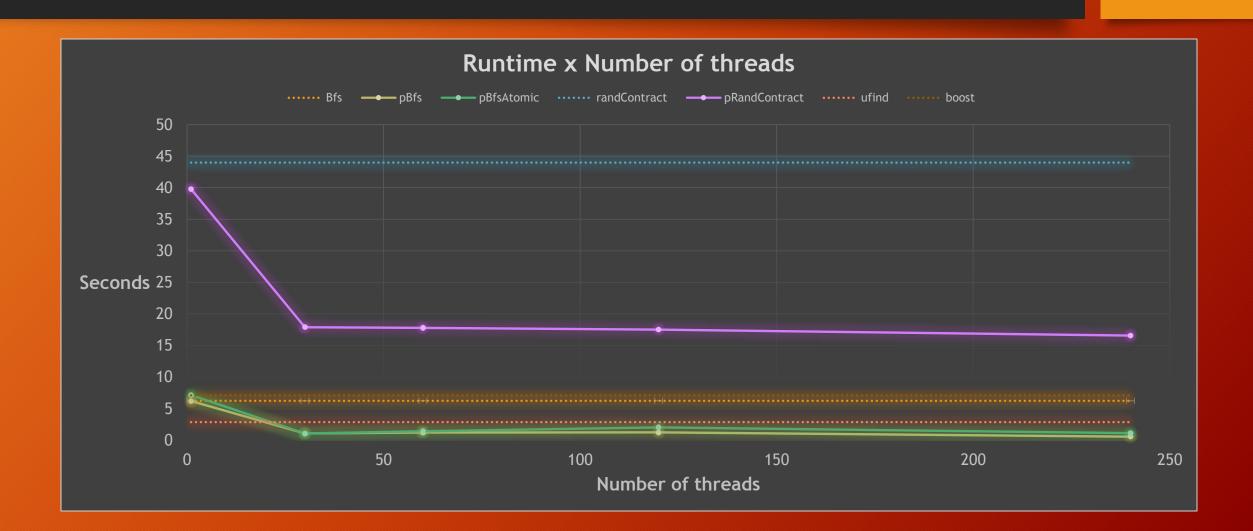
Memory usage varying the number of vertices



Runtime varying the number of components



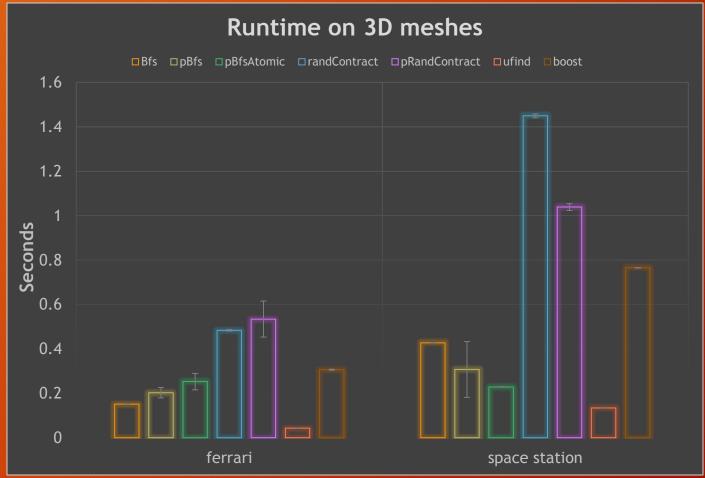
Runtime varying the number of threads



Runtime on real world graphs



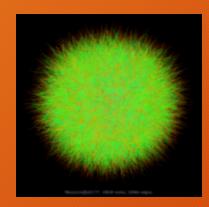
~350 thousand vertices ~700 thousand edges 753 components



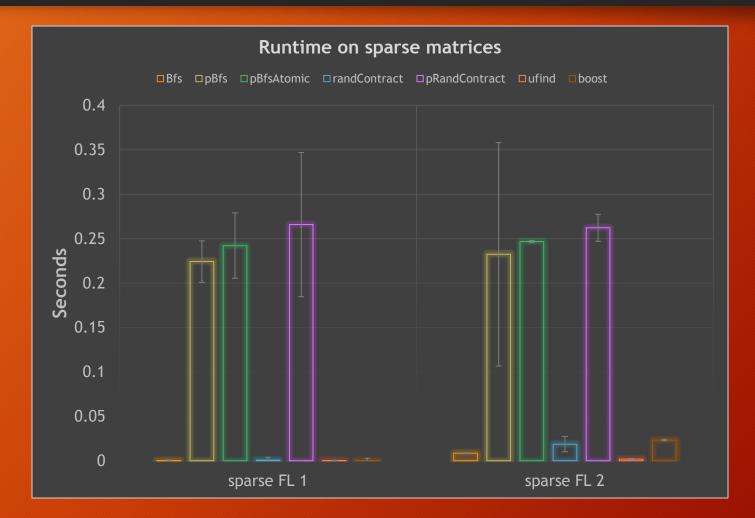


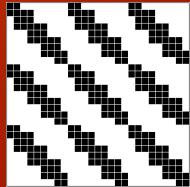
~1.2 million vertices ~2.4 million edges 6617 components

Runtime on real world graphs



608 vertices 1216 edges 1 component





~10 thousand vertices ~31 thousand edges 135 components

Conclusions and lessons learned

- Different algorithms are right for different situations
 - Best overall: union find (even though traversal has better complexity)
 - Best for very large graphs and plenty of cores: parallel traversal with atomics
- Should use algorithm that matches graph datastructure (not worth the time switching)
- Extra threads can help runtime up to a limit
- Boost is very bloated in terms of memory and not difficult to beat in runtime
- It is not a good idea to take measurements on a shared Xeon phi on the weekend before final presentation (5 hour queue)

Thank you