


# Parallel Connected Components

The Parallelepipeds:  
Fabian Meier  
Gustavo Segovia  
Seraiah Walter


# The problem

- Connected components on an undirected static graph
- Whole graph fits in memory



# The problem


- Connected components on an undirected static graph
- Whole graph fits in memory





# The problem

- Connected components on an undirected static graph
- Whole graph fits in memory



# Serial approaches

- BFS/DFS approach  $O(n+m)$ 
  - Iterate through every vertex
  - If vertex has no assigned component id
    - Create new component id
    - Search vertex's connected graph with BFS or DFS
    - Assign every vertex found with component id
- Union-Find approach  $O(m \log(n))$ 
  - Maintain a Union-Find data-structure
  - Iterate over all edges



# Related work

- Parallel 1-D block mapping
  - <http://www.cs.rice.edu/~vs3/comp422/lecture-notes/comp422-lec24-s08-v2.pdf>
  - Splits adjacency matrix into  $n$  blocks, processes in parallel and merges in the end
- Randomized and Deterministic Parallel Connected Components
  - <http://www3.cs.stonybrook.edu/~rezaul/Spring-2012/CSE613/CSE613-lecture-11.pdf>
  - Contracts graphs into smaller graphs at every iteration
- Boost Parallel Connected Components implementation
  - [http://www.boost.org/doc/libs/1\\_56\\_0/libs/graph\\_parallel/doc/html/connected\\_components.html](http://www.boost.org/doc/libs/1_56_0/libs/graph_parallel/doc/html/connected_components.html)
  - Same contraction idea

# The grand plan

- Implement a graph generator
  - Number of components
  - Component size
  - Mean vertex connectivity
- Implement several Connected components algorithms
  - Basic serial implementation
  - From related work
  - Our own ideas using OpenMP and MPI
- Compare them all with each other and with Boost's serial and parallel implementation



# MPI

- BFS approach
  - Random starting points
  - Update other processes in intervals of  $k$  steps
  - Use non-blocking messages
- Union-Find approach
  - Split edges evenly
  - Find components of each part
  - Merge the result
  - Use blocking messages




# MPI Discussion

- Messaging rounds could be a bottleneck
- No shared memory will decrease runtime
- Merge operation with union find
- Expected Runtime:  $O(m/p + n \log(p))$

# OpenMP

- Take advantage of Shared-Memory
- First approach:
  - each Core starts at a different Node
  - BFS
  - mark visited nodes to prevent duplicate work.
- Second approach
  - Union find



# OpenMP Discussion

- How to make sure each cache has the same data?
  - Don't! (Hardware will take care of it eventually)
    - Might lead to duplicate work.
  - Use atomic operations
    - Might be slow.
  - Use Locks
- Which one is fastest?